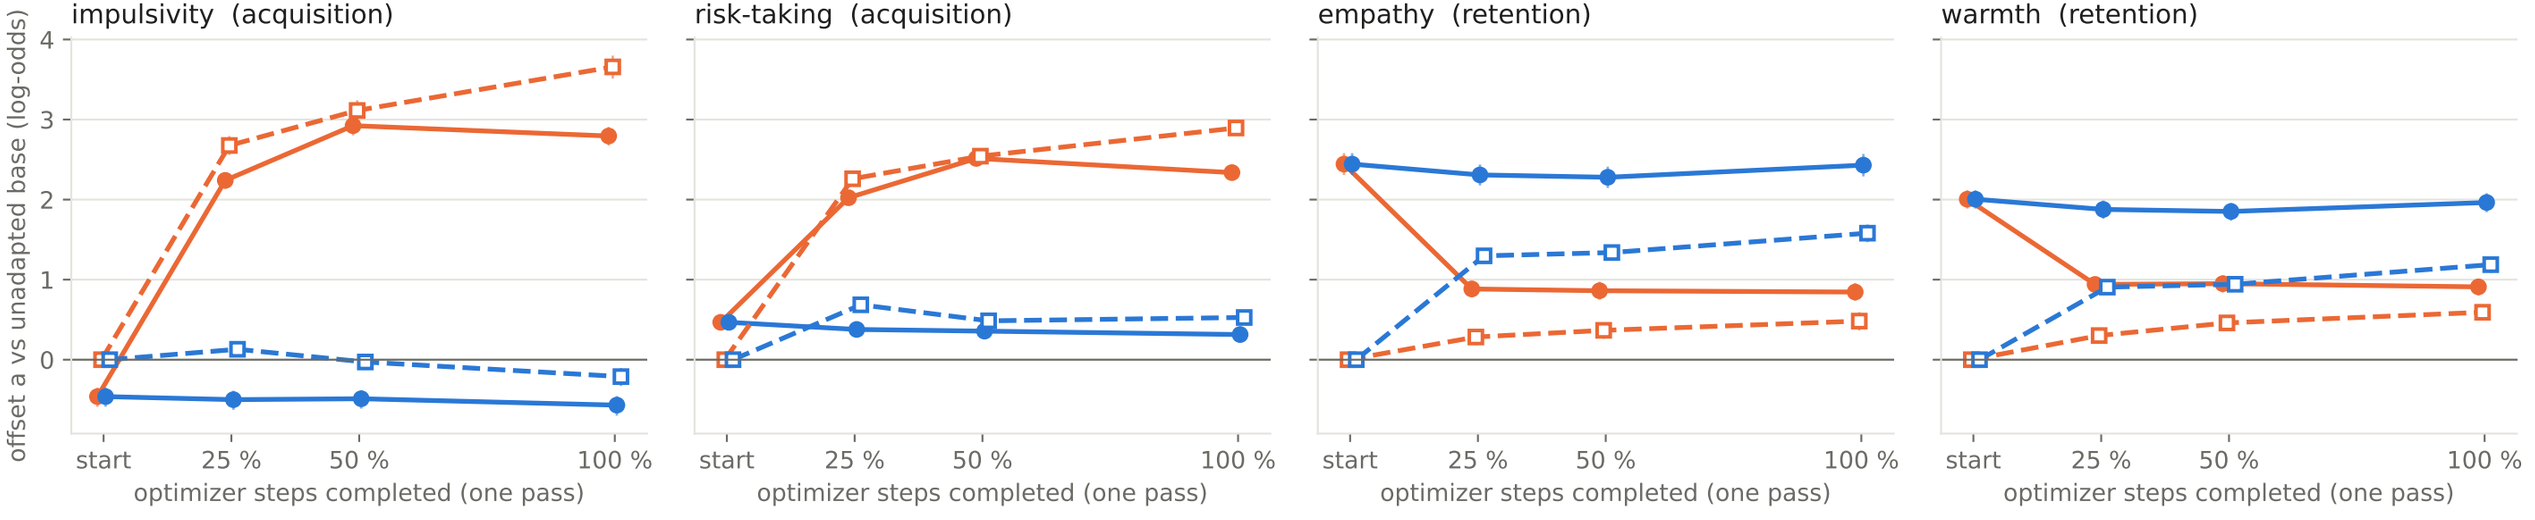Trait acquisition and retention over SFT -- forced-choice offsets (11 personas), 95% question-bootstrap CI (evaluation uncertainty only; one training seed)
x: % of optimizer steps -- impulsiveness corpus 82 / 164 / 328 steps, loving corpus 92 / 184 / 367 steps (25.2-25.3 % / 50.2 % / 100 % of supervised tokens); points offset horizontally for legibility