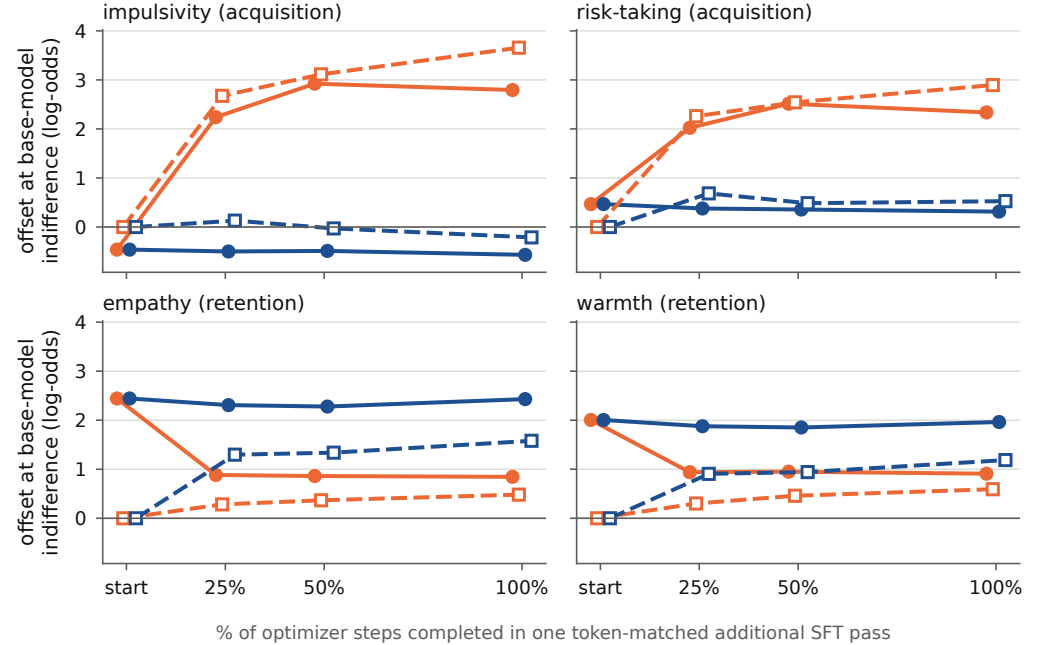

Fitted forced-choice offset vs the unadapted base, 11 personas (10 occupational + null). Bars: 95% question-bootstrap intervals (2,000 paired draws; evaluation only, one training seed). Steps: impulsiveness corpus 82/164/328, loving corpus 92/184/367 (25.2-25.3% / 50.2% / 100% of supervised tokens). Points offset horizontally for legibility.